

Parallel Programming with Parsl



Douglas Friedel (NCSA)

Complex Workflows on HPC Systems

- Can be challenging to manage
- May need to be partly re-written depending on what system you are on
- May not scale well depending on need
- Synchronization between different task is difficult
-
-
- Parsl aims to address these issues
(<https://parsl-project.org>)



What is Parsl?

- Flexible and scalable parallelization framework
- Purely written in Python
- Simple to use:
 - `@python_app` decorator for python functions
 - `@bash_app` for anything else
- Extensible from 1 to 8192 compute nodes, up to 0.25 M workers
- Built in interfaces for many systems (AWS, Slurm, Kubernetes, etc.)
- Can be used to implement many different parallel paradigms

What Can Parsl Do?

- Manage large number of parallel and serial tasks
- Synchronize input and output files for each task
 - Tasks with input files are held until all files are available
- Run your code on any system (laptop or HPC) with no changes
 - Just need to change your configuration

```
@python_app
def hello():
    return 'Hello World!'

print(hello().result())
```

Hello World!

```
@bash_app
def echo_hello(stdout='echo-hello.stdout'):
    return 'echo "Hello World!"'

echo_hello().result()


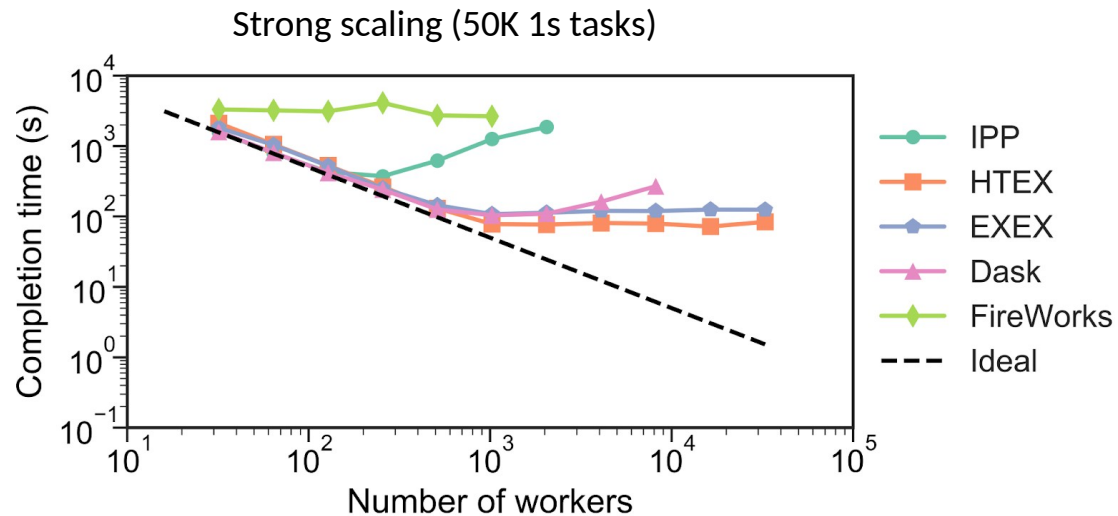
with open('echo-hello.stdout', 'r') as f:
    print(f.read())
```

Hello World!

Parsl Performance

HTEX and EXEX outperform other Python-based approaches

Framework	Maximum # of workers [†]	Maximum # of nodes [†]	Maximum tasks/second [‡]
Parsl-IPP	2048	64	330
Parsl-HTEX	65 536	2048 [*]	1181
Parsl-EXEX	262 144	8192 [*]	1176
FireWorks	1024	32	4
Dask distributed	4096	128	2617



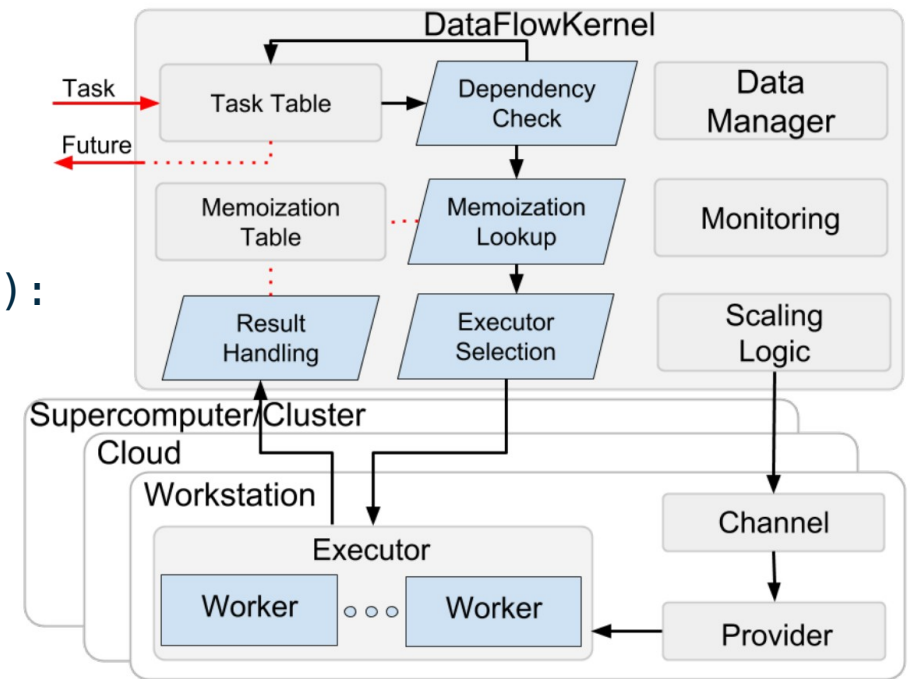
Babuji et.al. DOI:[10.1145/3307681.3325400](https://doi.org/10.1145/3307681.3325400)

How Do You Use Parsl?

- Basic use is to just add a decorator to your python functions

@python_app

@bash_app



How Do You Use Parsl?

- Call the decorated functions as normal

```
In [1]: value = random_multiply(5)
In [2]: value.result()
Out[2]: 4.934032990642622
```

```
In [1]: print(f"Return value is {echo('Hello world').result()}")
Out[1]: 0
In [2]: for line in open('std.out', 'r').readlines():
...:     print(line)
Out[2]: Hello world
```

Parsl and Files

- Any file passed as input or given as output needs to be wrapped in a Parsl File object
- Can handle local and remote files via
 - ftp/http/https
 - rsync
 - Globus file transfer

```
File("/local/scratch/myinput.txt")
```

```
File("https://github.com/Parsl/parsl/blob/master/README.rst")
```

```
File("globus://037f054a-15cf-11e8-b611-0ac6873fc732/unsorted.txt")
```


Parsl Structure

- Parsl Apps and Files are *futures*, a placeholder for an object which may not be available yet
- Any Python object (primitives, classes, etc.) can be passed as arguments
 - Must be serializable
- All `import` statements must be inside the function body
- `stdout/stderr` from remote machines may not be captured
-


Parsl Tutorial

- Fully hands on using Jupyter Notebooks (all you need is a browser and wifi connection)
- Parsl basics
 - Writing Apps
 - Parsl Blocks
- Configuring Parsl
- Handling files in Parsl
- Writing workflows
- Basic debugging



What is Parsl?

- Flexible and scalable parallelization framework
- Purely written in Python
- Simple to use:
 - `@python_app` decorator for python functions
 - `@bash_app` for anything else
- Each task can be run on a different resource
 - -



Questions?

This material is based in part upon work supported by the Department of Energy, National Nuclear Security Administration, under Award Number *DE-NA0003963*.